

Train/Validation I have created a 85:15 train/validation split, so that there are 581 points in the training set and 102 points in the validation set.

Network Architecture

- fully connected layer with 512 neurons and sigmoid activation
- fully connected layer with 512 neurons and sigmoid activation
- output layer with 2 neurons and softmax activation




Figure 1: Train and test accuracy over number of epochs.

Data	Label	Prediction
6 1 3 1 2 1 3 1 1	0	1
5 3 3 4 2 4 3 4 1	0	1

Table 1: Misclassified examples.

Performance The model is trained for 50 epochs. The train accuracy can reach 97.2% while the test accuracy can reach 97.0% (Figure 1). Training accuracy tends to be lower than test accuracy during experiments. I think that this may be due to lack of training data, or lack of features. Bias values tend to be around 0, but the weight values are not all 0 (Figure 2). Bias and weights for layer 2 have higher values.




Figure 2: Final weights and biases of trained model. Each weight and bias has 3 points associated with it: maximum, minimum, and mean. These values indicate the largest, smallest, and average of values in the matrix/vector.

Table 1 shows two misclassified points. The points in the Breast Cancer dataset are vectors describing patient's breasts, like clump thickness and cell shape. These attributes have values ranging from 1 to 10. The two examples are supposed to be labeled as malignant, but instead labeled as benign. I can see how they would be misclassified as benign since the attributes aren't too large in value.

Explanation The dataset does not contain images, so I have not used a CNN for this dataset. Instead, I use a feed-forward network with two fully-connected layers of 512 neurons. I have used a sigmoid activation function this time because I have found that classification accuracy increases from using sigmoid. This may be because the sigmoid curve does a good job of modeling the breast cancer data. Batch normalization and reducing the batch size also helps increase accuracy.

I have removed dropout from this network because I think that the network does not suffer from overfitting. In fact, I find that the network underfits because of lack in training data. The last layer also uses softmax function to output probability distribution over the classes. The loss function is binary cross-entropy because there are only two classes. This network only took about 5 minutes to train on a CPU.